

Evaluating History-Aware Design Skills for Coding Agents:

A Preregistered Matched-Pair Study of Memi 2.7.3

Memi Research and Design Engineering

Public technical disclosure and reproducibility report

1 August 2026

Abstract

This study does not establish that Memi is better. Coding agents can receive repository-specific design skills, but successful automation does not establish that the resulting interface is usable, that a skill transfers safely to a new repository, or that historical success should authorize future use. We evaluate the exact published Memi 2.7.3 artifact in a preregistered matched-pair study spanning web, Expo, and SwiftUI repositories. Thirty-six serialized agent cells formed 18 baseline/Memi pairs under the same provider, model, effort, task contract, and verification rules. All 36 receipts passed content-addressed admission. Fourteen cells were excluded only from rendered grading, leaving 10 complete quality pairs across two tasks. Three blinded model graders scored a 100-point interface rubric. The mean Memi-minus-baseline score difference was +1.4 for Buzzr (one-sided 95% lower bound +0.2) and -0.4 for Paraform (lower bound -3.4); both bounds exceeded the preregistered -5 point non-inferiority margin. Neither task established superiority, and no quality estimate was possible for Nate. Across 21 task-by-resource tests, no comparison remained significant after Holm correction. The evidence motivated a fail-closed router released in Memi 2.7.4: history is isolated by exact task, repository, harness, skill identity, and skill-content hash; a harmful event suppresses immediately; recovery requires three later healthy prospective pairs. A chronological replay left both observed routes suppressed and repository-only. We release the executed notebook, receipts, exclusions, grader records, figures, policy replay, and software-channel provenance. The findings are model-graded product research, not independent practitioner validation.

Keywords: coding agents; design engineering; matched-pair evaluation; non-inferiority; LLM-as-a-judge; evidence-based routing; reproducible systems research

1 Introduction

Repository-level coding benchmarks have made execution in real codebases a central unit of evaluation [1]. Interactive-agent work similarly emphasizes outcome-based validation in realistic environments [2]. Frontend engineering adds a second problem: a patch may compile and pass a functional test while still degrading accessibility, hierarchy, responsiveness, or design-system consistency. Design-to-code benchmarks have documented the distance between visual generation and robust frontend implementation [3], but less is known about agents that receive reusable design instructions inside an existing codebase.

Memi is a design-skill and audit layer for coding agents. Its router can select repository-relevant skill content before an agent edits code. This creates a deployment question that is both statistical and operational: when does past evidence justify enabling a skill route, and what should happen after a quality or efficiency regression? A permissive history mechanism can leak evidence across repositories, models, or revised skill content.

A purely functional benchmark can miss interface regressions. A release checklist can verify distribution without answering either research question.

This paper reports V15, a prospective confirmatory study of the exact Memi 2.7.3 npm artifact, and the fail-closed policy that followed in 2.7.4. It asks:

- RQ1: Interface quality.** Within each frozen task with complete rendered evidence, is Memi’s model-graded design quality non-inferior to the baseline agent under a -5 point margin?
- RQ2: Function and resources.** How do functional acceptance, critical defects, token use, wall time, tool calls, retries, and provider failures differ within matched pairs?
- RQ3: Safe reuse.** Can exact-route historical evidence be replayed without look-ahead and converted into a policy that suppresses harmful routes without borrowing evidence from another task, repository, model, or skill revision?

The contributions are: (i) a preregistered, multi-platform matched-pair study; (ii) an evidence model that separates execution validity, functionality, renderability, and blinded quality; (iii) a chronological backtest of an exact-match, fail-closed router; and (iv) a public artifact package that binds the analysis to published software channels. The paper makes no general superiority or dollar-savings claim.

2 Related work and positioning

Repository-level agent evaluation. SWE-bench evaluates patches against real repository tests, demonstrating the importance of pinned code state and executable outcomes [1]. WebArena extends realistic, functional evaluation to interactive web tasks [2]. V15 adopts the same high-level principle—evaluate the produced artifact in its repository—but adds matched conditions, interface-quality grading, content hashes, and release provenance.

Frontend and design evaluation. Design2Code combines rendering-based metrics with human evaluation and reports persistent layout and visual-element failures in automated frontend generation [3]. V15 instead studies bounded edits to existing products. The outcome therefore includes task correctness, accessibility, design-system fit, responsiveness, and code quality rather than screenshot similarity alone.

Model-based grading. Strong language models can approximate preference judgments, but position, verbosity, and self-enhancement biases remain material [4]. We therefore blind condition labels, use three graders, report disagreement and ICC, exclude non-renderable evidence, and describe the result as *model-graded*. We do not substitute this panel for independent practitioner evidence.

Table 1. Frozen task strata and required behavior.

Fixture	Stack	Required change	Pairs
Buzzr	Expo	Accessible unread-count badge	6
Paraform	Web	Keyboard command menu	6
Nate	SwiftUI	Truthful reduced-motion state	6

3 Study design

3.1 Candidate, conditions, and tasks

The candidate was the immutable npm tarball for `@memi-design/cli@2.7.3`, SHA-256 `b2b56189f45b69a3cec41220ae3e9d50d6ecfb682aee5a64c1fe0ce0d0ec3e1e`. The baseline condition used the provider agent under the frozen task contract without a Memi-routed design skill. The Memi condition used the same contract and could receive at most one routed skill, recorded in its route receipt. Provider (Codex), model (`gpt-5.6-luna`), reasoning effort (low), maximum input tokens (450,000), verification rules, and fixture revision were held constant within each pair.

Three tasks were selected to exercise different interface stacks and forms of interaction (Table 1). Six pairs were run per task. Condition order was counterbalanced with frozen seed 2715. Cells executed serially; Xcode and Simulator received exclusive machine access.

3.2 Prospective execution and evidence admission

Each cell began from a clean clone at a pinned revision. Preparation, routing, patch, tool activity, usage, environment, isolated verification, and final run receipt were recorded. The admission pipeline rejected missing artifacts, duplicate trial identities, fixture mutation, hash disagreement, condition mismatch, or non-isolated verification. All evidence was content-addressed.

Figure 1 shows the central evidence boundary. Functional and resource outcomes use every admitted receipt. Rendered quality uses only cells with reviewable interface evidence; a grading-only exclusion does not erase the cell’s other outcomes. This distinction prevents render failures from being silently imputed while retaining observed execution failures as data.

3.3 Blinded interface-quality panel

Three model graders independently scored anonymized evidence. Condition labels were hidden. The rubric allocated 25 points to task and interaction correctness, 20 to accessibility, 15 each to visual hierarchy, design-system consistency, and responsive/adaptive behavior, and 10 to implementation quality. Dimension and total scores were aggregated by panel median. We report absolute disagreement and ICC(2,1), following the two-way random-effects reliability formulation of Shrout and Fleiss [7].

The rendering boundary is important. Buzzr’s evidence is a React Native testing renderer tree rather than native pixels. Paraform has registered desktop browser states, while its phone breakpoint intentionally replaces the desktop workspace. All 12 Nate cells lacked admissible rendered grading evidence. These constraints limit visual inference even when functional receipts are valid.

3.4 Evidence lineage

The earlier V14 study remains labeled 2.7.2 feasibility evidence: 12 cells and six pairs across Buzzr and Paraform, without Nate or a blinded rendered panel. It is neither relabeled nor pooled into V15. Release reconstruction found that 2.7.3 changed release metadata, guidance, and adapter-test stabilization rather than the previously benchmarked runtime paths; V15 nevertheless executed the exact 2.7.3 tarball instead of inferring behavior from that source comparison.

4 Statistical analysis

The experimental unit is a matched baseline/Memi pair within one task. For quality, the paired estimand is

$$\Delta_q = q_{\text{Memi}} - q_{\text{baseline}}. \quad (1)$$

We report the task-level mean and median, the exact sign test, a 10,000-sample paired bootstrap interval [5], and leave-one-pair-out sensitivity. The preregistered affirmative rule is non-inferiority, not superiority: the one-sided 95% lower bound must be strictly greater than -5 points on the 100-point rubric. The margin was frozen before execution.

Secondary outcomes are functional acceptance, critical defects, provider-reported input/output/reasoning tokens, wall time, tool calls, retries, and provider failures. Raw differences are Memi minus baseline. Tests are oriented so that higher functional acceptance and lower values elsewhere are favorable. Holm’s sequential correction is applied across task-level secondary tests [6]. Because tasks differ in platform and behavior, pooled cross-fixture results are descriptive. No token measurement is converted into dollars because billing records were not collected.

Missing quality pairs are not imputed. A complete functional receipt remains in secondary analysis even if rendered grading is unavailable. Critical-defect analysis requires complete blinded/rendered evidence and therefore has a smaller denominator. The executed notebook is the numerical source of record; the paper imports generated tables and figure assets from that pipeline.

5 Results

5.1 What V15 establishes—and what it does not

The short answer is deliberately narrow. V15 supports one bounded statement: on the two tasks with complete rendered evidence, the data rule out a design-quality loss of five or more rubric points under the preregistered model-panel procedure. It does not establish that Memi is better, generally faster, or cheaper. The third task has no admissible quality estimate, the quality panel contains only 10 complete pairs, none of the resource tests survived multiplicity correction, and billing observations were not collected.

5.2 All cells were admitted, but quality evidence covered only two tasks

All 36 frozen cells had unique, manifest-verified receipts. Fourteen exclusions applied only to rendered grading: all 12 Nate cells, Buzzr repeat 6 Memi, and Paraform repeat 5 baseline. Twenty-two cells received grades. Twenty formed 10 complete pairs—five Buzzr and five Paraform—while two graded cells remained unpaired after sibling exclusion. No value was imputed.

Prospective matched-pair design and evidence boundary

A renderability exclusion removes only the blinded visual score; it does not erase functional or resource evidence.




Figure 1. Prospective execution and evidence admission. All 36 cells entered functional and resource analysis. Fourteen prespecified, cell-level exclusions applied only to rendered grading; 22 cells were graded and 20 formed 10 complete quality pairs.

What the V15 study actually establishes

The result is a bounded safety finding, not a general claim that Memi produces better designs.




Figure 2. Plain-language claim boundary for V15. Clearing the orange -5 point margin means that a large decline was not observed on the two gradable tasks; it is not evidence that Memi outperforms the baseline.

5.3 RQ1: The study rules out a large decline, not a performance win

Figure 3 separates observed pair differences from the task-level decision statistic. The orange -5 line is an acceptable-loss boundary, not a threshold for winning: a decision bound to its right passes non-inferiority, while a mean to the right of zero still does not prove better performance. Buzzr’s mean difference was +1.4 points (median +1; two-sided 95% bootstrap interval 0.0 to +2.8); its one-sided lower bound was +0.2. Paraform’s mean was -0.4 (median +2; interval -4.2 to +2.8); its one-sided lower bound was -3.4 . Both lower bounds lie above the -5 margin. Exact one-sided sign-test p values were 0.1875 and 0.5, respectively; these do not establish superiority.

Table 2. Task-level blinded design-quality results generated from the executed analysis.

Task	Pairs	Mean Δ	Median Δ	One-sided 95% lower	NI gate
buzzr-tab-unread-badge	5	1.4	1	0.2	pass
paraform-command-menu	5	-0.4	2	-3.4	pass

The panel’s pooled ICC(2,1) was 0.531536. Pair-level abso-

lute disagreement varied across tasks and conditions, so the median score should not be read as a human consensus estimate. No blinded critical defect was recorded in the 10 complete graded pairs. Nate has no quality estimate; functional success cannot fill that missing construct.

5.4 RQ2: No secondary test survived multiplicity correction

Functional acceptance differed by -0.1667 for Buzzr, 0 for Nate, and $+0.1667$ for Paraform. All Holm-adjusted p values were 1. Critical-defect differences were zero where renderable pairs existed; Nate had no admissible critical-defect estimate.

Figure 4 reports the four continuous resource outcomes. Four of 21 task-by-resource bootstrap intervals excluded zero before multiplicity correction, but none of the 21 task-level secondary tests rejected after Holm correction. The estimates are heterogeneous: for example, Nate showed wider token and wall-time intervals than the other tasks. The data do not support a general efficiency claim, and provider-reported tokens are not billing cost.

Sparse count outcomes are retained in the machine-readable

Quality result: no large decline observed

Better performance was not established: this is a narrow non-inferiority finding on two tasks, not a general design-quality win.




Figure 3. Blinded design-quality evidence. Dots are complete pairs, pale spans are two-sided 95% paired-bootstrap intervals, solid dots are task means, and vertical ticks are the one-sided decision bounds. The orange reference is the -5 point acceptable-loss margin; the zero line denotes equality. Passing the orange line means “no large decline observed,” not “Memi is better.”

secondary table: tool calls, retries, and provider failures. Nate’s provider-failure interval favored Memi before correction, while Buzzr had a small positive retry difference; neither authorized a corrected claim. Serial execution controls local contention but does not estimate concurrent throughput.

6 Fail-closed evidence routing

6.1 Policy derived from exact route identity

The 2.7.4 policy defines a route by the exact tuple

$$R = (\text{task class, repository fingerprint, provider, model, effort, skill ID, skill content hash}). \quad (2)$$

Evidence from another tuple is ineligible. Fitness evidence v2 binds blinded quality data and its SHA-256 to that identity while preserving v1 loading. A v1 negative event may suppress; automation-only v1 evidence cannot promote or recover. This asymmetry makes legacy uncertainty conservative.

Figure 5 expresses the state machine. A quality regression, or a joint catastrophic regression of at least 50% more tokens and 25% more latency, suppresses immediately. Normal execution then falls back to repository-only discovery. Recovery is possible only through three later healthy, prospective, exact-match v2 pairs. New harmful evidence resets recovery progress. A route without exact evidence abstains rather than borrowing a favorable record.

This policy is a safety mechanism, not evidence that a skill improves a future frontend. Its purpose is to prevent weak, stale, or mismatched evidence from turning a skill into a default route.

6.2 RQ3: Replay confirmed policy mechanics, not product efficacy

The chronological backtest replayed 12 store-write-eligible events in pair-complete order: 10 blinded-quality v2 events and two automation-only negative v1 events. Seven additional

rows were chronology-only. At each cutoff, the replay consumed only events with timestamps at or before that cutoff.

Paraform/**design-extract** was suppressed by its first exact event. Buzzr/**atomic-design** remained healthy through four events and was suppressed at 09:11 UTC after a quality regression. Neither route accumulated three eligible recovery pairs. Nate abstained in all six repeats; therefore no Nate route state or suppression claim exists.

The replay validates deterministic policy mechanics and temporal isolation. It does not show that this event stream represents future repositories, models, providers, or skill revisions. The detailed chronological replay remains in the reproducibility package because it validates timing and isolation, rather than a user-facing quality or efficiency outcome.

7 Frontend and release validation

7.1 Independent product-surface audit

The public website was audited from clean checkouts, excluding deployment metadata, build output, generated community examples, and other non-product sources from design review. The controlled baseline reproduced a serious Lighthouse **label-content-name-mismatch** finding in all nine reports across three routes. A footer accessible-name correction aligned the computed name with visible text. All nine post-fix reports contained no blocking accessibility finding, and the unchanged Lighthouse assertion gate moved from exit 1 to exit 0. The Chromium light-mode suite passed 55/55 tests before and after. The selected axe rules also passed in both conditions because the experimental Lighthouse rule was outside that frozen selection.

This before/after result supports one narrow conclusion: the observed accessible- name defect was removed under the unchanged gate. Local performance variation is not interpreted causally, and local Astro execution is not deployment proof.

Paired resource differences with 95% bootstrap intervals

Six matched pairs per task. Intervals reflect provider-reported usage and serialized wall time, not billing cost.




Figure 4. Task-level paired resource differences. Points are mean Memi-minus- baseline differences and lines are two-sided 95% paired-bootstrap intervals. Each task contributes six pairs. Negative values favor lower Memi resource use; units are shown per panel.

7.2 Release verification as disclosure, not efficacy evidence

Memi 2.7.4 shipped the exact-route policy through trusted-publisher OI DC. Independent channel receipts verified npm latest, a nine-cell Node 20/22/24 install matrix on Linux, macOS, and Windows, the GitHub tag and platform binaries, Action v2, Homebrew stable, GHCR latest, the MCP Registry, and the production website artifact. The public gate returned failures: [] and parityEligible: true. GHCR was verified by an authenticated promotion workflow; anonymous pullability is not claimed.

These receipts establish that the disclosed implementation and paper were published consistently. They are not observations in the V15 efficacy study and do not increase its sample size.

8 Discussion

RQ1. Within the two tasks that produced complete rendered pairs, the lower confidence bounds exceeded the preregistered non-inferiority margin. This is evidence against a quality loss larger than five rubric points under the observed conditions. It is not evidence that Memi is superior, and it says nothing about Nate’s rendered quality.

RQ2. Functional and resource estimates did not yield a corrected secondary claim. Wide and task-dependent intervals are scientifically useful: they show that one aggregate efficiency number would hide platform-specific behavior. A future cost study must collect billing data directly.

RQ3. Exact-route replay changed the operational decision

even where task-level non-inferiority passed. This is not contradictory. The router’s event-level quality regression rule is deliberately more conservative than the paper’s task-level mean non-inferiority statement. Research inference asks what the observed paired sample supports; deployment policy asks whether any exact route has accumulated enough safe prospective evidence to remain enabled.

The broader lesson is methodological: compile success, test success, rendered quality, policy eligibility, and distribution parity are different constructs. A credible agent-system report should preserve those boundaries instead of turning them into one composite score.

9 Threats to validity

Internal validity. Pairing, counterbalancing, frozen fixtures, serial execution, and isolated verification reduce known confounding. They do not eliminate stochastic provider behavior. Two orchestration deviations are recorded: the admission harness was separated from the exact 2.7.3 runtime, and a stalled partial attempt was retained as retry evidence while only the later complete receipt supplied the trial outcome.

Construct validity. The 100-point rubric covers important interface dimensions but is not a validated human-preference instrument. Model graders may share biases with the evaluated agent family. Buzzr lacks native pixel evidence; Paraform’s mobile state is intentionally a desktop-only placeholder; Nate lacks rendered grades entirely. Functional acceptance and critical defects use different admissible denominators.

Statistical conclusion validity. Quality inference uses five complete pairs per estimable task, and secondary analyses use

Fail-closed routing: unproven skills do not become defaults

A route needs exact evidence to run. Harm turns it off immediately; only later exact-match evidence can turn it back on.




Figure 5. Fail-closed skill-fitness policy. Exact identity is checked before historical evidence can affect execution. Suppression is immediate; promotion from legacy automation evidence is prohibited; recovery requires three later healthy prospective pairs on the same route.

six. Bootstrap intervals and leave-one-pair-out checks expose instability but cannot create power. The -5 margin is a preregistered engineering tolerance, not a universally accepted clinical or perceptual threshold. Exact sign tests and Holm correction are conservative at this sample size.

External validity. The study covers three repositories, one provider/ model/effort configuration, one Apple-silicon workstation, and one day of execution. It does not estimate multi-user throughput, maintainability, independent practitioner preference, billing cost, or performance on unregistered repositories.

Researcher independence. The study was executed and reported by the team developing Memi. The prospective freeze, machine-readable ledgers, condition blinding, immutable artifacts, and public negative results improve auditability but do not replace independent replication.

10 Reproducibility and disclosure

The source of record is the executed `analysis.ipynb`. The public package contains the frozen protocol and plan, environment, rubric, exact candidate provenance, task contracts, 36 evidence receipts, exclusion and deviation ledgers, blinded grader records, pair-level and task-level tables, 10,000-sample bootstrap results, leave-one-pair-out sensitivity, figures, fitness-policy inputs and replay, LaTeX source, PDF, and checksum inventories. The report builder rejects stale derived files, duplicate JSON keys, non-finite numbers, digest mismatch, and disagreement between release channels.

The exact candidate source commit is `2e29c5c656ac34242369eac9840838619ad113e1`. The published 2.7.4 implementation is bound to `8aa4649f412bbcaaf2af4ee209bf79016566f035`. V14 remains separately labeled and hashed. All headline numbers can be recomputed without the prose layer.

Conflict-of-interest statement. Memi Research and Design Engineering developed the evaluated system and conducted this study. This document is a public technical disclosure, not a peer-reviewed or independently commissioned evaluation.

Human-participant and data statement. No human participants, private user records, or billing observations were collected. Interface-quality labels were produced by three blinded model graders. Repository fixtures and generated agent artifacts are identified by pinned revisions and content hashes.

Falsification criteria. An affected claim must be withdrawn if any candidate checksum, frozen trial identity, fixture revision, evidence hash, exclusion record, notebook output, one-sided lower bound, or public-channel binding fails independent recomputation. Condition leakage, look-ahead replay, an unreported exclusion, or a critical accessibility or security finding also invalidates the corresponding claim.

11 Conclusion

In this preregistered study, Memi 2.7.3 met a -5 point model-graded quality non-inferiority criterion on two renderable tasks, while a third task lacked the evidence required for a quality claim. Secondary functional and resource comparisons did not survive multiplicity correction. The resulting 2.7.4 policy treats historical design-skill evidence as route-specific, suppresses on harm, and requires prospective exact-match recovery. The contribution is not a claim that Memi is generally better; it is a reproducible method for separating what an agent executed, what reviewers could judge, what statistics support, and what a production router should be allowed to do.

What must change before a better, faster, or cheaper claim

The next study must use native rendered evidence on every platform, a preregistered independent human practitioner panel, measured billing records, and a larger confirmatory sample sized before execution. It must separately test quality superiority and efficiency, rather than treating a non-inferiority pass or provider token count as either result. The accompanying 2.7.5 plan specifies the harness, skill, router, and release gates required before those claims may be made; new evidence must remain isolated by exact route identity and skill-content hash.

References

- [1] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. R. Narasimhan, "SWE-bench: Can Language Models Resolve Real-World GitHub Issues?," in *International Conference on Learning Representations*, 2024.

- [2] S. Zhou, F. F. Xu, H. Zhu, X. Zhou, R. Lo, A. Sridhar, X. Cheng, T. Ou, Y. Bisk, D. Fried, U. Alon, and G. Neubig, “WebArena: A Realistic Web Environment for Building Autonomous Agents,” in *International Conference on Learning Representations*, 2024.
- [3] C. Si, Y. Zhang, R. Li, Z. Yang, R. Liu, and D. Yang, “Design2Code: Benchmarking Multimodal Code Generation for Automated Front-End Engineering,” arXiv:2403.03163, 2024.
- [4] L. Zheng, W.-L. Chiang, Y. Sheng, et al., “Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena,” in *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2023.
- [5] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. Chapman & Hall, 1993.
- [6] S. Holm, “A Simple Sequentially Rejective Multiple Test Procedure,” *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979.
- [7] P. E. Shrout and J. L. Fleiss, “Intraclass Correlations: Uses in Assessing Rater Reliability,” *Psychological Bulletin*, vol. 86, no. 2, pp. 420–428, 1979.